ECE 46300 - Fall 2024

Transport Layer III

Vishal Shrivastav



Transmission Control Protocol (TCP)

Reliable, In-order Delivery to Applications

- Application Multiplexing
- ACKs and Retransmissions
- Flow Control
- Congestion Control

Byte Stream Abstraction



Connection-oriented

- Pairwise < sender, receiver > connection is established before sending data
- Used to maintain connection-specific state, e.g., initial seq num, window scaling factor, window size, etc.

Transmission Control Protocol (TCP)

Reliable, In-order Delivery to Applications

- Application Multiplexing
- ACKs and Retransmissions
- Flow Control
- Congestion Control

Byte Stream Abstraction

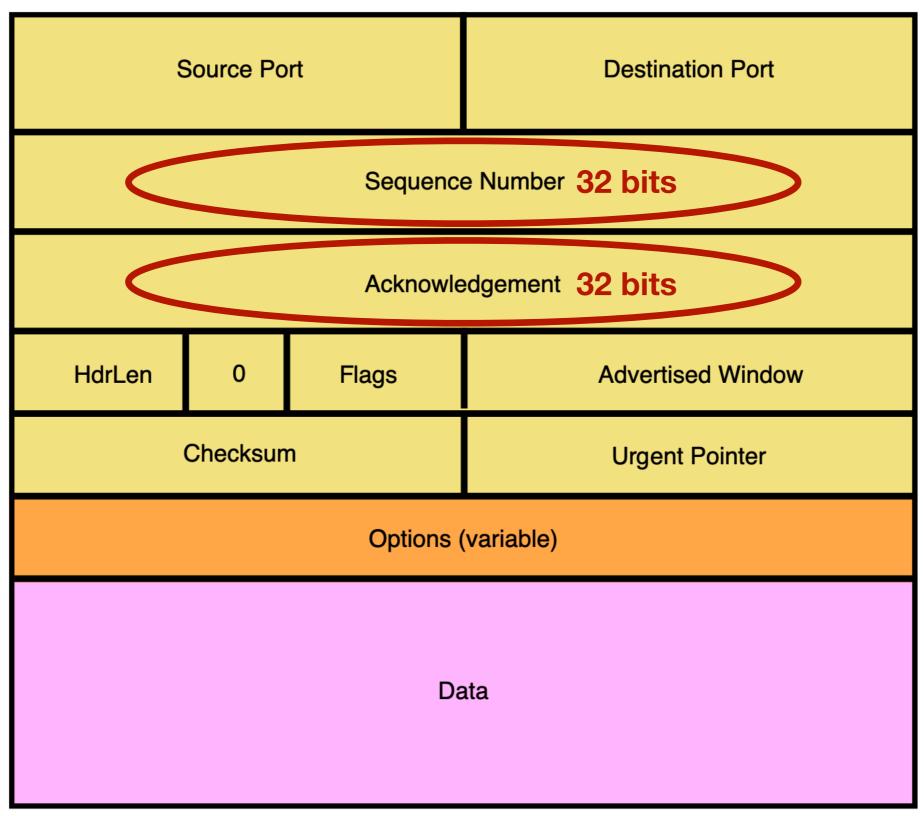
Communicates with the applications using a byte stream and <u>not</u> packets

Connection-oriented

- Pairwise < sender, receiver > connection is established before sending data
- Used to maintain connection-specific state, e.g., initial seq num, window scaling factor, window size, etc.

ACKs

Sequence and Acknowledgement Num

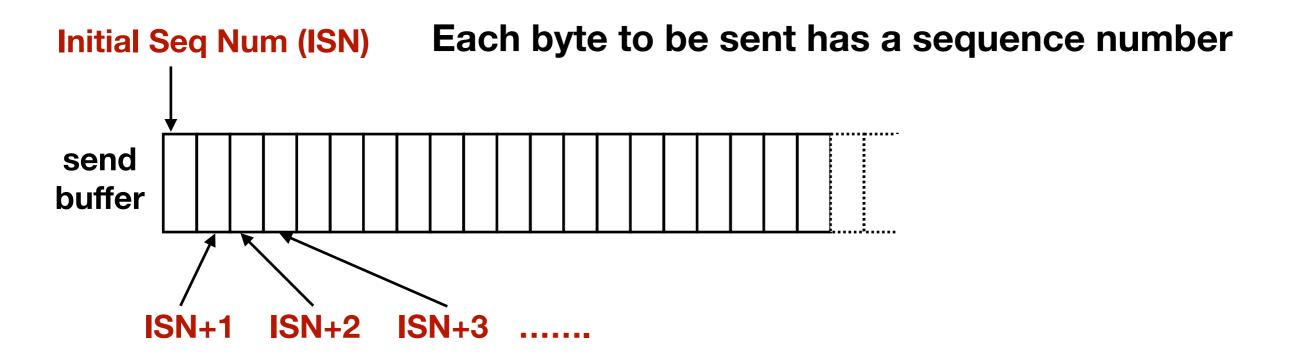


Previously we focused on <u>packets</u>

- Sequence number referred to packet number
- ACKs contained received sequence/packet numbers
- Window size was expressed in terms of # of packets

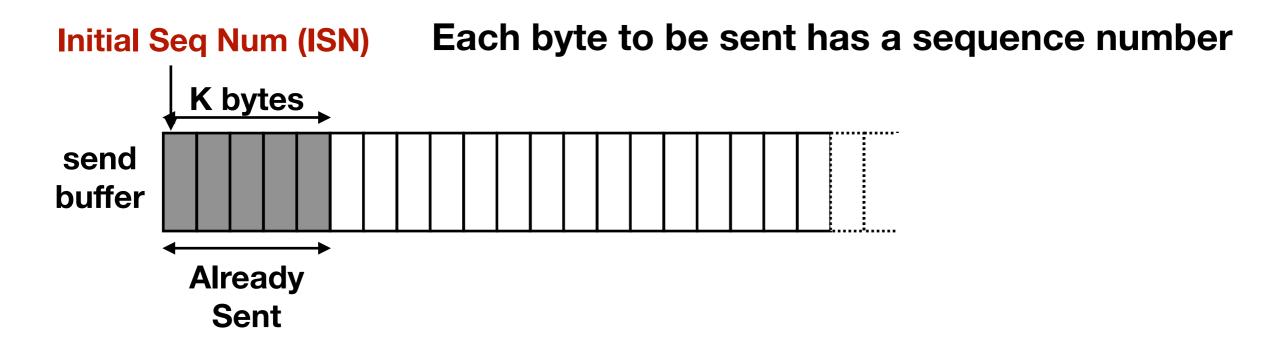
But, TCP focuses on <u>bytes</u>

- Sequence number refers to byte number of bytes being sent
- ACKs contain received sequence/byte numbers
- Window size expressed in terms of # of bytes

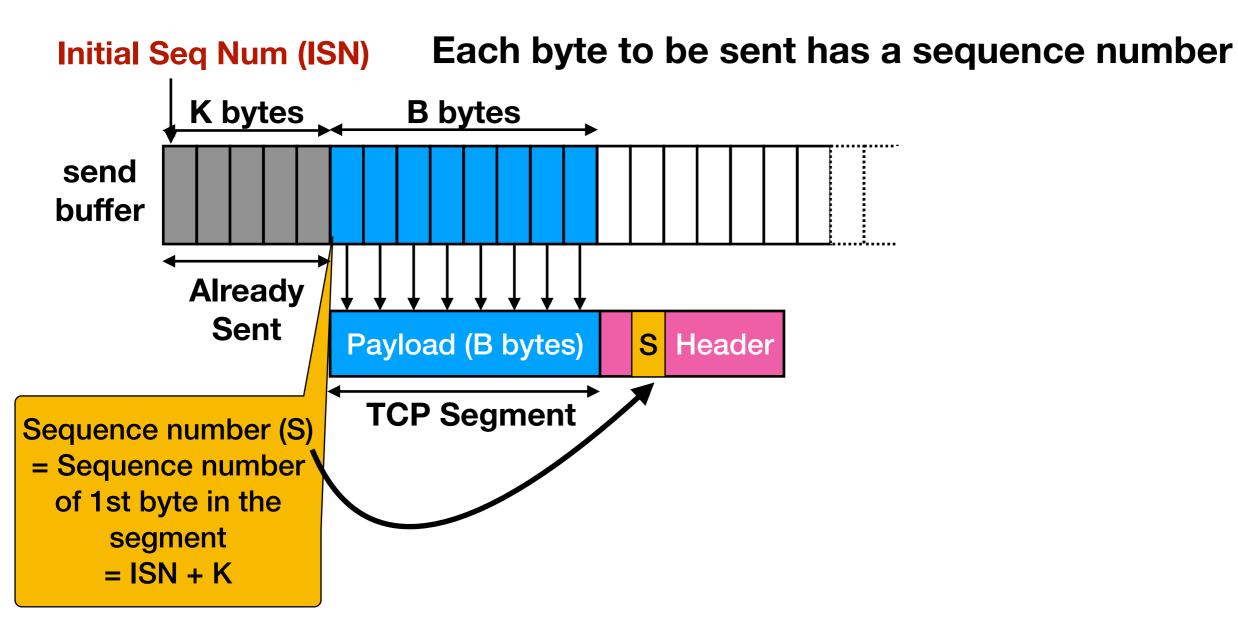


Suppose sequence number is represented using k bits, What will be the sequence number of $(2^k + 1)^{th}$ byte?

Answer: Same as the sequence number of 1st byte, i.e., sequence numbers wrap around!

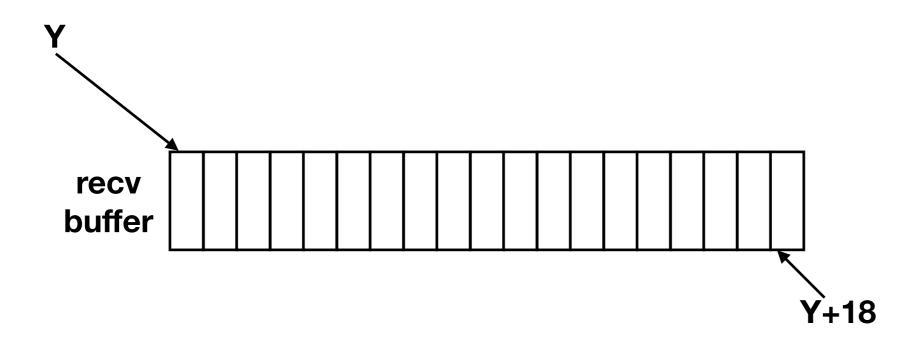


Initial Seq Num (ISN) Each byte to be sent has a sequence number send buffer Already Payload (B bytes) S Header TCP Segment

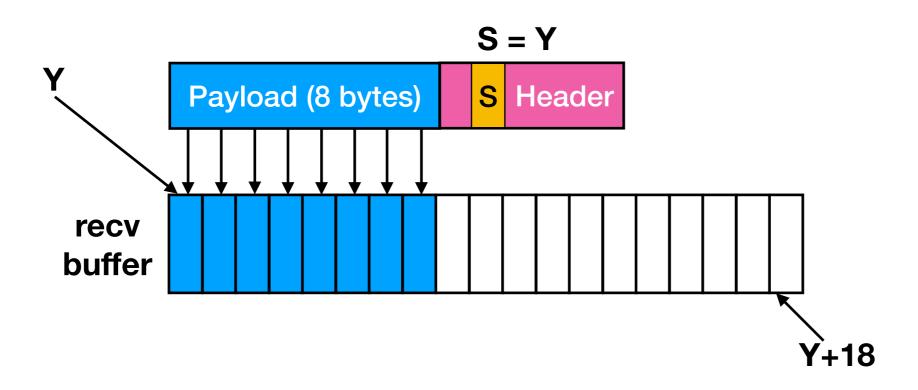


- Sender sends TCP packets with TCP segments (byte stream)
 - Sequence number field in the packet header equals the sequence number of 1st byte in TCP segment (= S)
 - TCP segment contains B bytes, with sequence numbers
 - S, S+1, S+2,, S+B-1

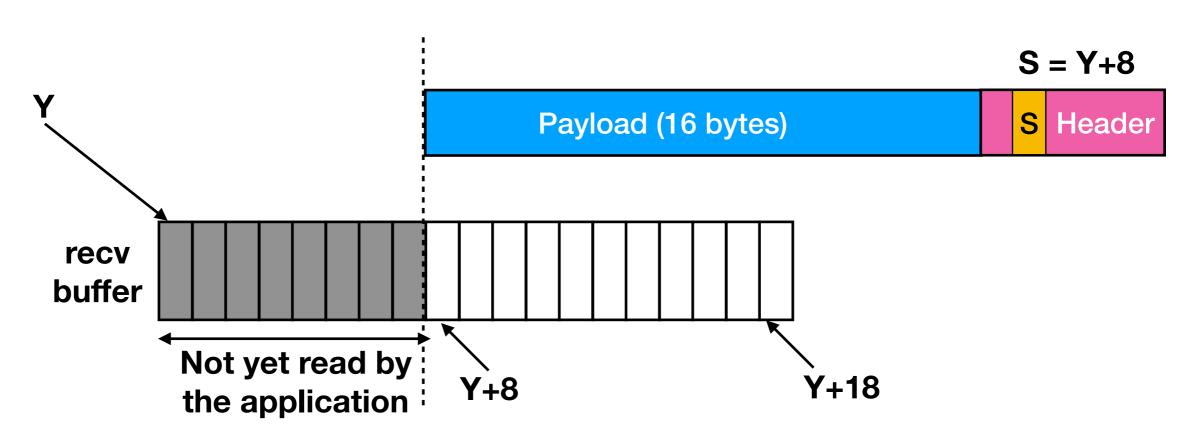
- Suppose application has read all bytes until sequence number Y-1
 - Receive buffer can hold bytes with sequence numbers starting from Y to Y + len(recv buffer) - 1



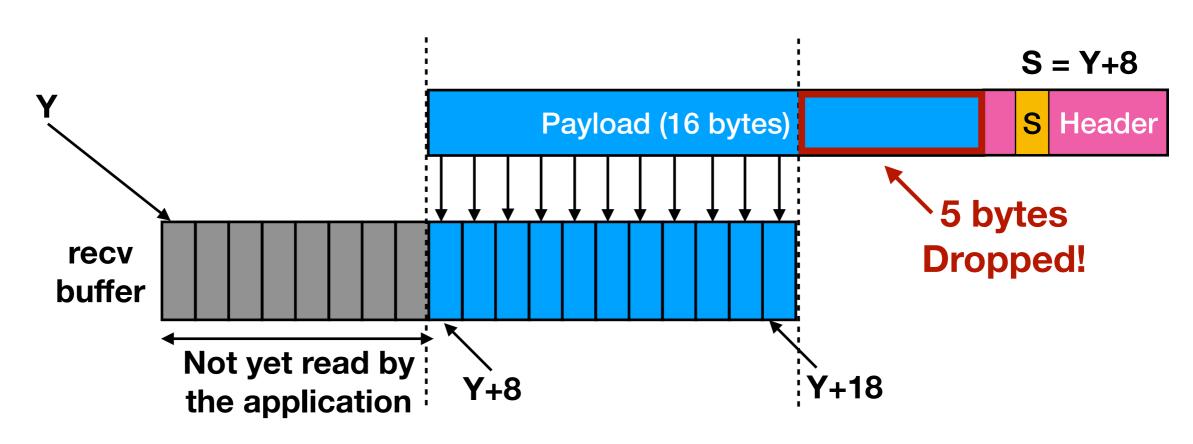
- Suppose application has read all bytes until sequence number Y-1
 - Receive buffer can hold bytes with sequence numbers starting from Y to Y + len(recv buffer) - 1



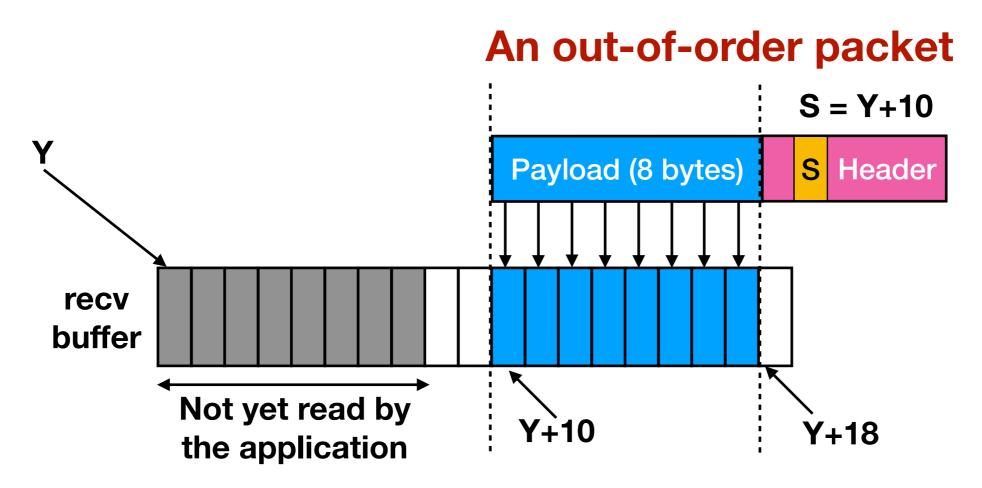
- Suppose application has read all bytes until sequence number Y-1
 - Receive buffer can hold bytes with sequence numbers starting from Y to Y + len(recv buffer) - 1



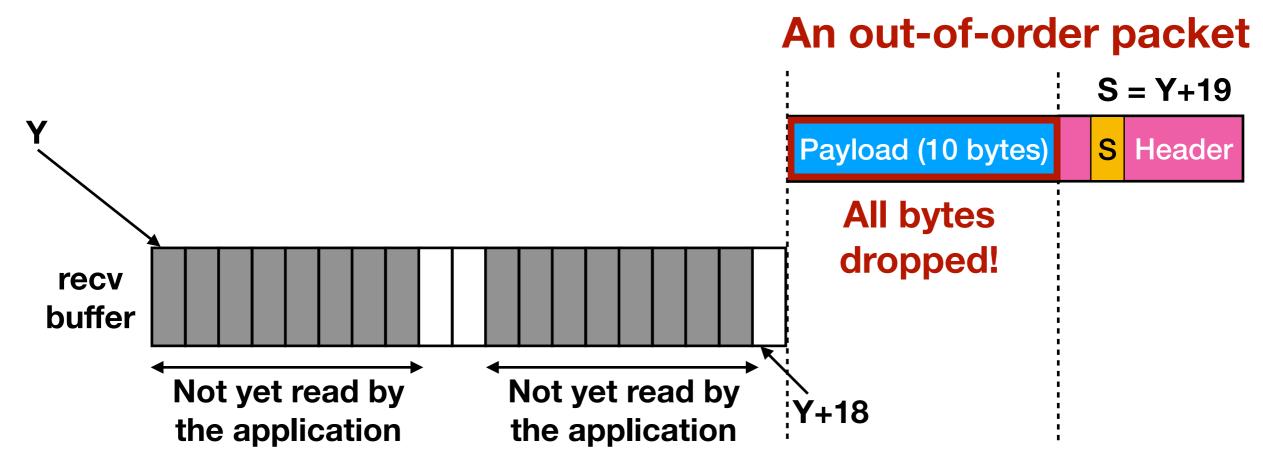
- Suppose application has read all bytes until sequence number Y-1
 - Receive buffer can hold bytes with sequence numbers starting from Y to Y + len(recv buffer) - 1



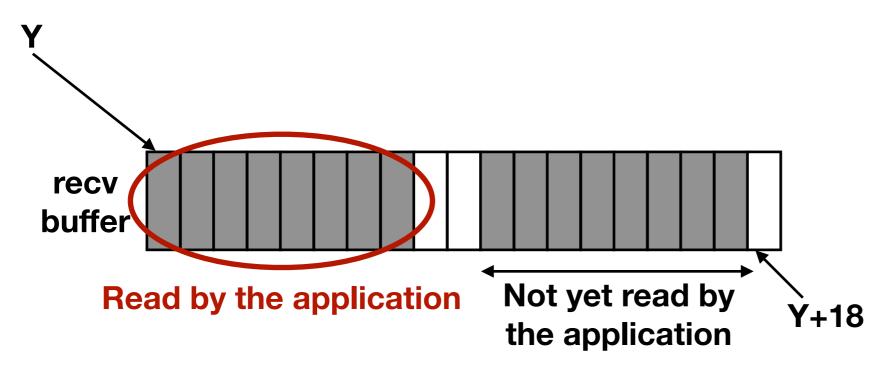
- Suppose application has read all bytes until sequence number Y-1
 - Receive buffer can hold bytes with sequence numbers starting from Y to Y + len(recv buffer) - 1



- Suppose application has read all bytes until sequence number Y-1
 - Receive buffer can hold bytes with sequence numbers starting from Y to Y + len(recv buffer) - 1



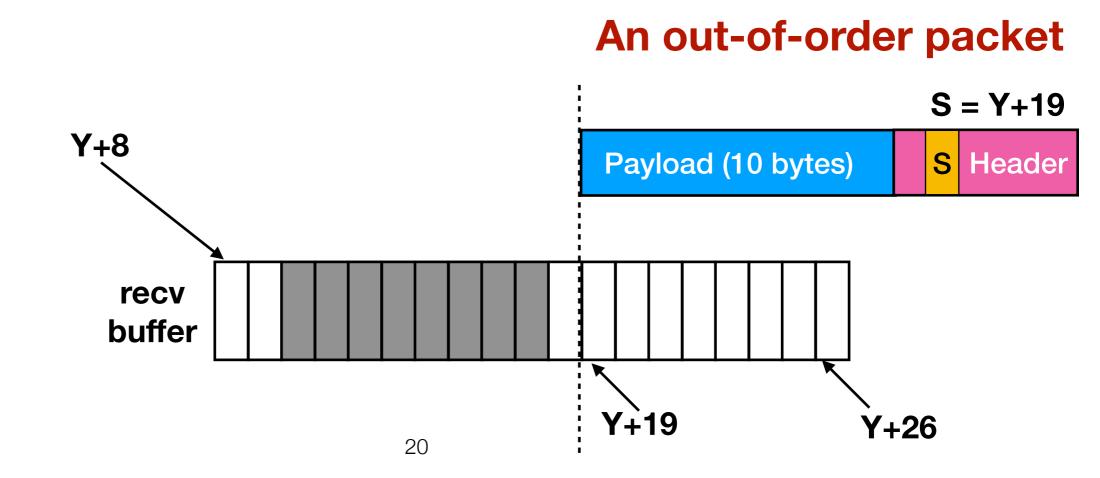
Suppose application has read all bytes until sequence number Y+7



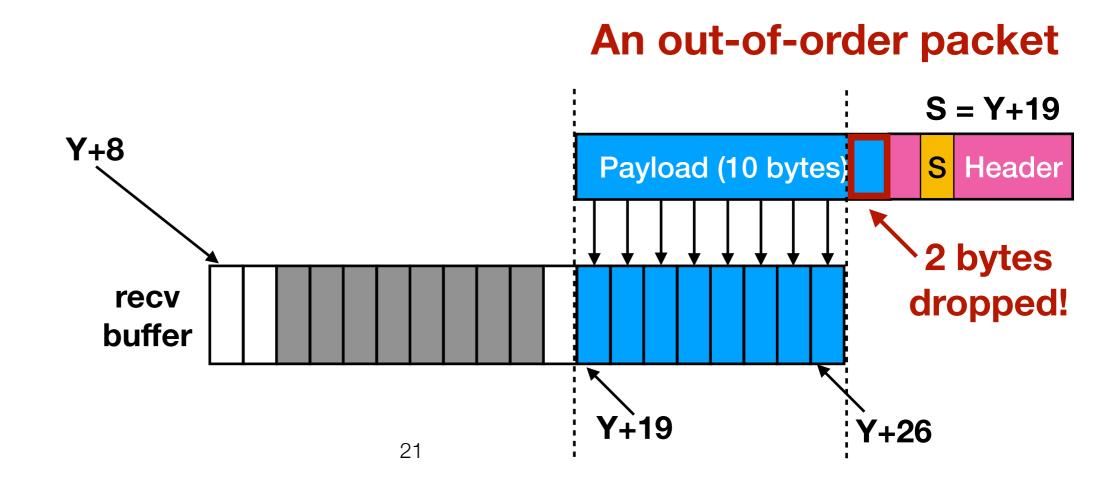
- Suppose application has read all bytes until sequence number Y+7
 - Receive buffer can now hold bytes with sequence numbers starting from Y+8 to Y+8 + len(recv buffer) - 1



- Suppose application has read all bytes until sequence number Y+7
 - Receive buffer can now hold bytes with sequence numbers starting from Y+8 to Y+8 + len(recv buffer) - 1



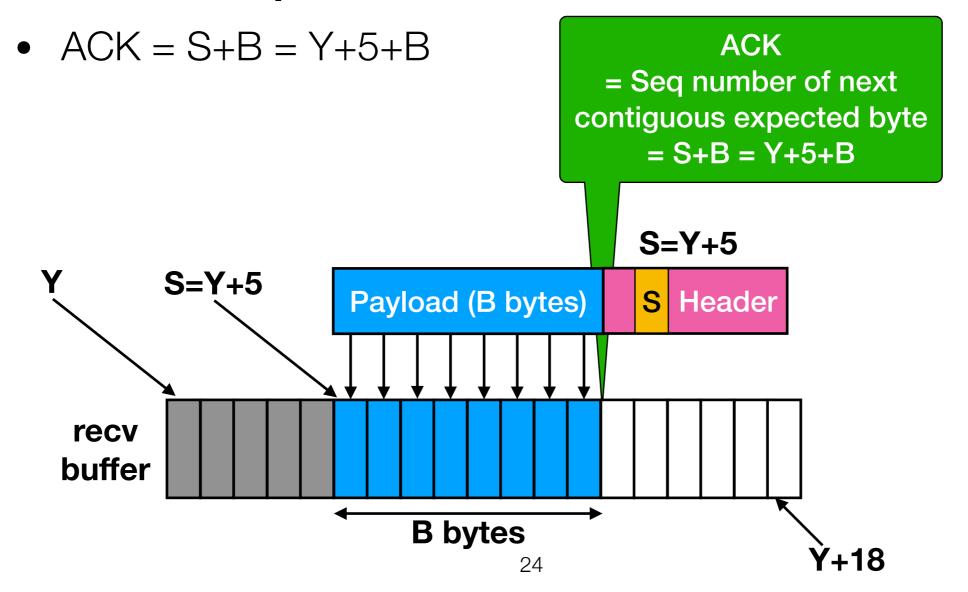
- Suppose application has read all bytes until sequence number Y+7
 - Receive buffer can now hold bytes with sequence numbers starting from Y+8 to Y+8 + len(recv buffer) - 1



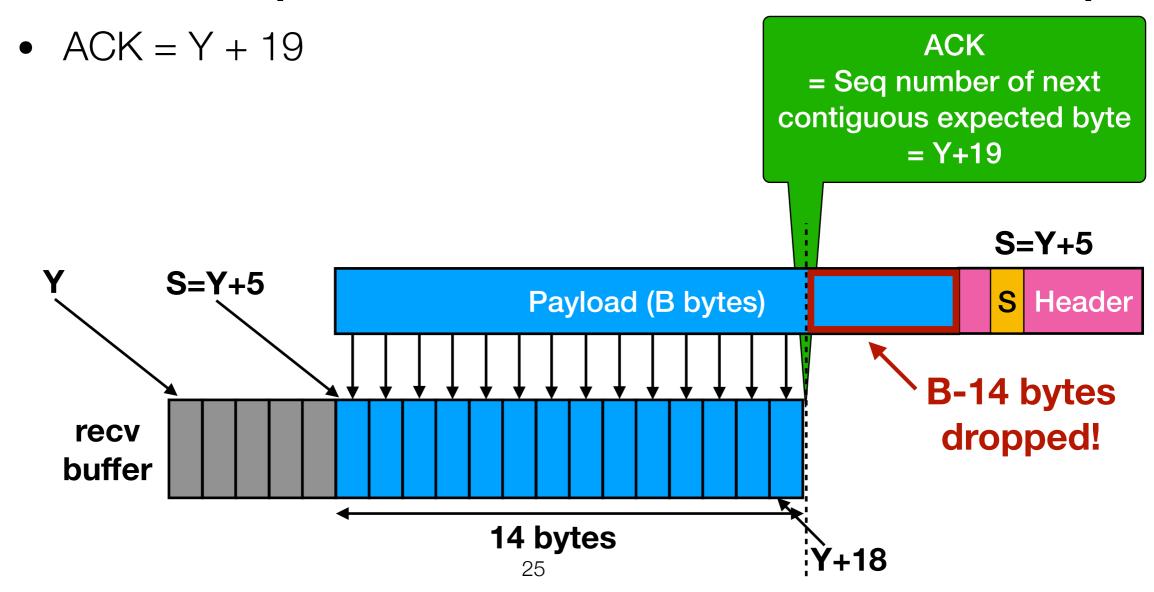
- Sender sends TCP packets with TCP segments (byte stream)
 - Sequence number field in the packet header equals the sequence number of 1st byte in TCP segment (= S)
 - TCP segment contains B bytes, with sequence numbers
 - S, S+1, S+2,, S+B-1
- Upon receipt of a TCP packet, receiver sends an ACK
 - Most popular TCP implementations use Cumulative ACK

- Sender sends TCP packets with TCP segments (byte stream)
 - Sequence number field in the packet header equals the sequence number of 1st byte in TCP segment (= S)
 - TCP segment contains B bytes, with sequence numbers
 - S, S+1, S+2,, S+B-1
- Upon receipt of a TCP packet, receiver sends an ACK
 - Most popular TCP implementations use Cumulative ACK
 - ACK contains the sequence number of next contiguous byte expected by the receiver

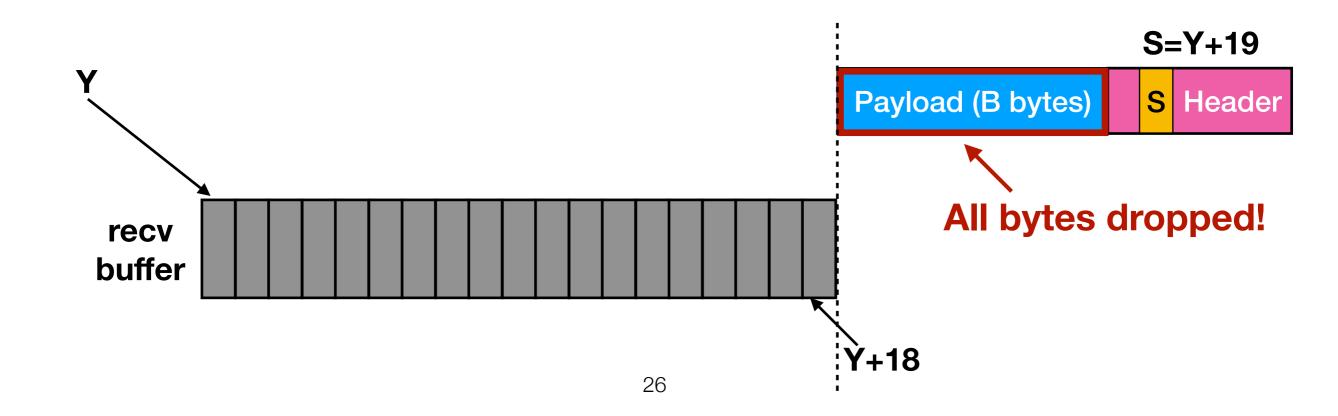
- Suppose bytes upto Y+4 received and bytes upto Y-1 read by application
 - On receiving a TCP packet with sequence num S and B bytes of data,
 - Case 1a: All bytes prior to S have been received, i.e., S = Y+5, and received packet fits in the available recv buffer space



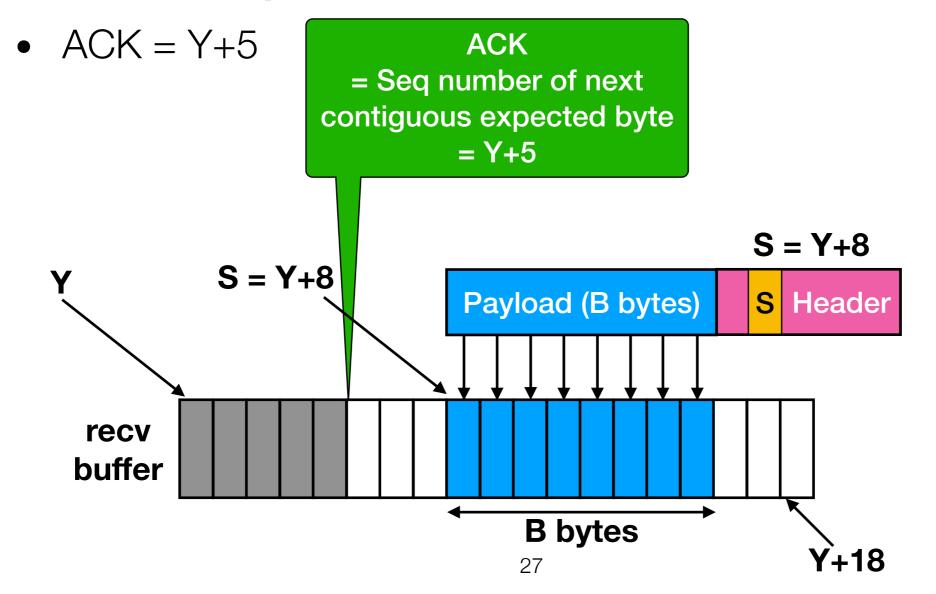
- Suppose bytes upto Y+4 received and bytes upto Y-1 read by application
 - On receiving a TCP packet with sequence num S and B bytes of data,
 - Case 1b: All bytes prior to S have been received, i.e., S = Y+5, and received packet doesn't fit in available recv buffer space



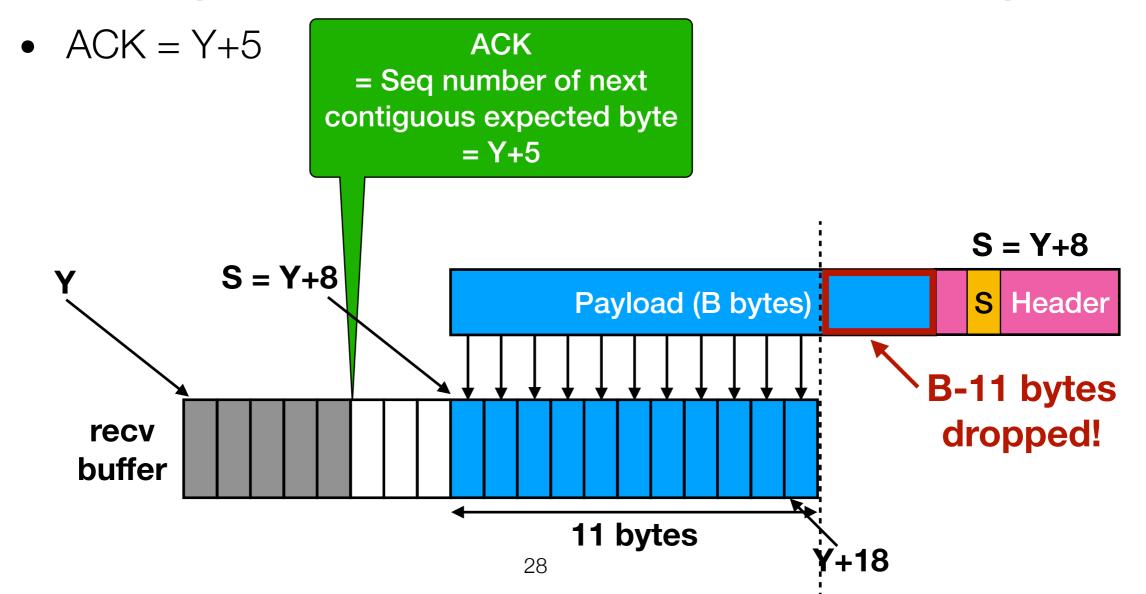
- Suppose bytes upto Y+4 received and bytes upto Y-1 read by application
 - On receiving a TCP packet with sequence num S and B bytes of data,
 - Case 1c: All bytes prior to S have been received, i.e., S = Y+19, and recv buffer is full
 - Entire packet is discarded ... No bytes received ... No ACK sent!



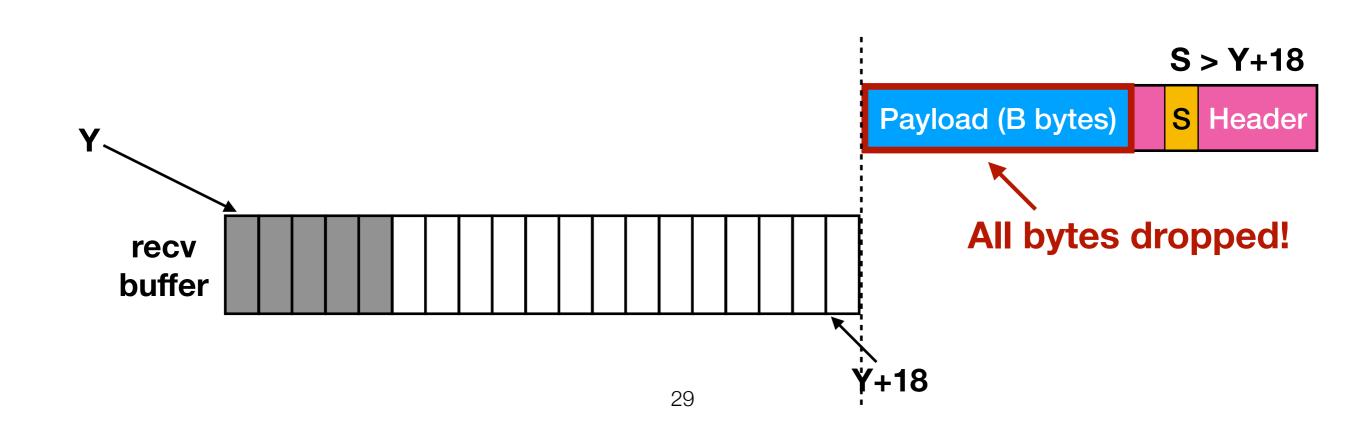
- Suppose bytes upto Y+4 received and bytes upto Y-1 read by application
 - On receiving a TCP packet with sequence num S and B bytes of data,
 - Case 2a: Received an out-of-order packet, i.e., S > Y+5, and the received packet fits in available recv buffer space



- Suppose bytes upto Y+4 received and bytes upto Y-1 read by application
 - On receiving a TCP packet with sequence num S and B bytes of data,
 - Case 2b: Received an out-of-order packet, i.e., S > Y+5, and received packet doesn't fit in available recv buffer space



- Suppose bytes upto Y+4 received and bytes upto Y-1 read by application
 - On receiving a TCP packet with sequence num S and B bytes of data,
 - Case 2c: Seq num of out-of-order pkt exceeds recv buffer size
 - Entire packet is discarded ... No bytes received ... No ACK sent!



Retransmissions

TCP Retransmission Mechanism

- Under what conditions should a sender retransmit a packet?
 - Assuming Cumulative ACKing,
 - Duplicate ACKs
 - Timeout
- What packets should a sender retransmit?
 - Assuming Cumulative ACKing, use Go-Back-N

TCP Retransmission Mechanism

- Under what conditions should a sender retransmit a packet?
 - Assuming Cumulative ACKing,
 - Duplicate ACKs
 - Timeout
- What packets should a sender retransmit?
 - Assuming Cumulative ACKing, use Go-Back-N

Duplicate ACKs (Fast Retransmit)

Example:

- Sender sends 100B packets with ISN 100
- Sequence numbers of sent packets:
 - 100, 200, 300, 400, 500, 600, 700, 800, 900,
- Assume 5th packet (sequence num 500) is lost but no others
- Stream of ACKs will be:
 - 200, 300, 400, 500, 500, 500, 500, 500, 500,

Stream of duplicate ACKs; implies 500 has not been delivered TCP retransmits after receiving 3 Duplicate ACKs

Why not retransmit immediately after receiving the first Duplicate ACK?

500 could have been **reordered/delayed** in the network, and **not** lost... Waiting for 3 Dup ACKs before retransmitting can avoid unnecessary retransmissions

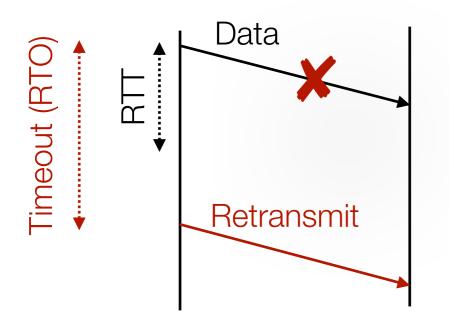
Why need timeouts? Why (Duplicate) ACKs are not enough?

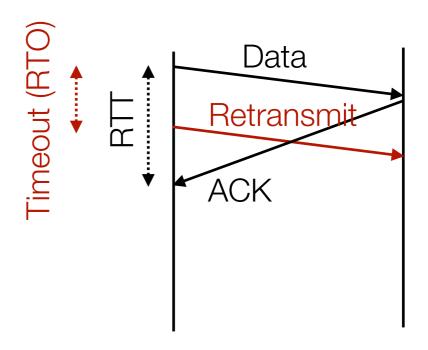
- Why need timeouts? Why (Duplicate) ACKs are not enough?
 - ACKs can only be generated if <u>some</u> packets are being received
 - What if a TCP stream only has 1 packet, and that packet is lost?
 - Or, the last packet in a TCP stream is lost?
 - Or, a retransmitted packet is lost?
 - ... will need a timer to detect loss!

- Why need timeouts? Why (Duplicate) ACKs are not enough?
 - ACKs can only be generated if <u>some</u> packets are being received
 - What if a TCP stream only has 1 packet, and that packet is lost?
 - Or, the last packet in a TCP stream is lost?
 - Or, a retransmitted packet is lost?
 - ... will need a timer to detect loss!
- Most TCP implementations have a single timer
- TCP resets timer whenever new data is ACKed
- Retransmits when the timer expires
- **RTO** (Retransmit Time Out) value inside the TCP code is used to configure the timeout value

36

How to set the value of RTO?





Too large => waiting too long to retransmit

Too small => unnecessarily retransmitting

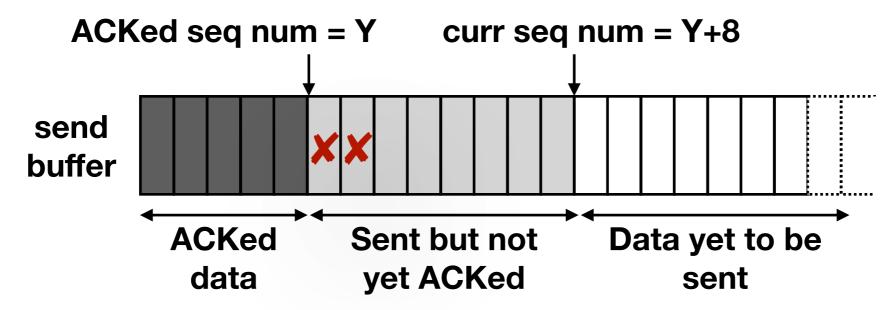
- Ideally, RTO = RTT ... but need a very accurate estimation of RTT
 - Very hard to do in practice, because of variable delays such as queuing delay
- Most TCP implementations use a conservative value for RTO, e.g., 500ms
 - Hence, retransmission by timeout is very slow!
 - So, TCP relies on Duplicate ACKs for "Fast Retransmit"

TCP Retransmission Mechanism

- Under what conditions should a sender retransmit a packet?
 - Assuming Cumulative ACKing,
 - Duplicate ACKs
 - Timeout
- What packets should a sender retransmit?
 - Assuming Cumulative ACKing, use Go-Back-N

Go-Back-N

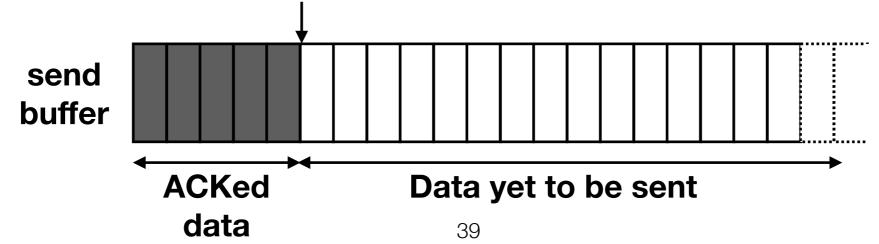
ACKed seq num = highest contiguous sequence num ACKed by receiver **curr seq num** = sequence num of next byte that sender would send



Suppose some bytes starting from Y are lost → Will receive Dup ACKs with value Y

On receiving 3 Dup ACKs, sender will do Go-Back-N

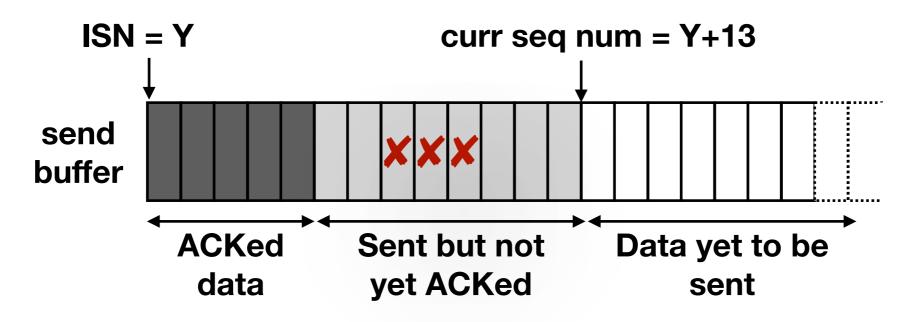
curr seq num = ACKed seq num = Y — start (re)transmitting from byte Y



Other TCP ACK and Retransmission Mechanisms

- Most TCP implementations use Cumulative ACK with Go-Back-N
- However, some TCP implementations also use
 - Selective ACK with Selective Retransmit
 - May result in better performance, e.g., fewer retransmitted packets
 - ... at the cost of higher ACKing overhead
 - Selective ACK vs. Cumulative ACK
 - ... and higher implementation complexity
 - Selective Retransmit vs. Go-Back-N

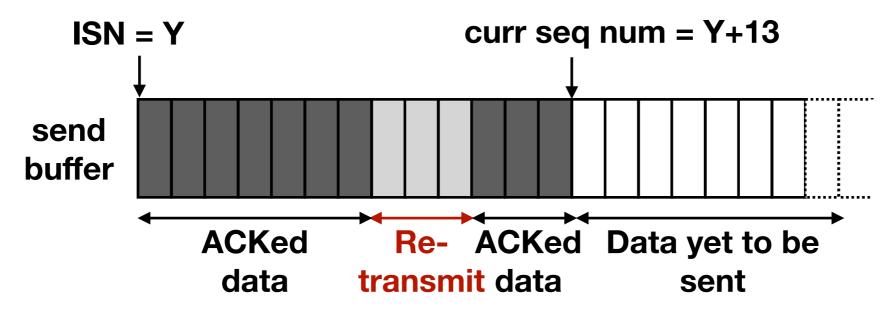
Selective ACK w/ Selective Retransmit



Suppose bytes Y+7, Y+8, Y+9 are lost

Will receive **Selective ACK** = (up to Y+6, plus Y+10, Y+11, Y+12)

On receiving the Selective ACK, sender will do Selective Retransmit



Transmission Control Protocol (TCP)

Reliable, In-order Delivery to Applications

- Application Multiplexing
- ACKs and Retransmissions
- Flow Control
- Congestion Control

Byte Stream Abstraction

Communicates with the applications using a byte stream and <u>not</u> packets

Connection-oriented

- Pairwise < sender, receiver > connection is established before sending data
- Used to maintain connection-specific state, e.g., initial seq num, window scaling factor, window size, etc.

Questions?