ECE 46300 - Fall 2024

Transport Layer IV

Vishal Shrivastav



Transmission Control Protocol (TCP)

Reliable, In-order Delivery to Applications

- Application Multiplexing
- ACKs and Retransmissions
- Flow Control
- Congestion Control

Byte Stream Abstraction

Communicates with the applications using a byte stream and <u>not</u> packets

Connection-oriented

- Pairwise <sender, receiver> connection is established before sending data
- Used to maintain connection-specific state, e.g., initial seq num, window scaling factor, window size, etc.

Transmission Control Protocol (TCP)

Reliable, In-order Delivery to Applications

- Application Multiplexing
- ACKs and Retransmissions
- Flow Control
- Congestion Control

Byte Stream Abstraction

Communicates with the applications using a byte stream and <u>not</u> packets

Connection-oriented

- Pairwise < sender, receiver > connection is established before sending data
- Used to maintain connection-specific state, e.g., initial seq num, window scaling factor, window size, etc.

TCP Connection Set up

- A TCP connection is identified using 5-tuple
 - <source IP, destination IP, source Port, destination Port, Protocol=TCP>
- Connection set up used to exchange initial sequence number (ISN)
 - Also used to exchange window scaling factor (used in flow control)
- A receiver needs to know the ISN of the sender to figure which sequence number marks the start of the byte stream being received
- Each endpoint choses its ISN randomly
 - Why endpoints not just choose ISN as 0 (or some fixed number)?

Why random ISN?

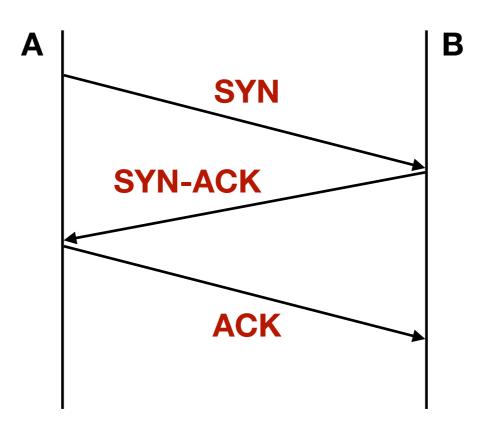
Practical Issue 1: Port numbers are re-used

- Suppose a TCP connection between A and B is closed
- ... and a new TCP connection is opened between A and B with the same source and destination port numbers as the previous connection
- But suppose a packet with sequence num 100 from previous TCP connection was still in-flight
- B will think packet belongs to new TCP connection if ISN of new connection
 was 0 ... less chance of this happening if ISN is chosen randomly
- e.g., if random ISN for new TCP connection was >100, then not an issue in the above example

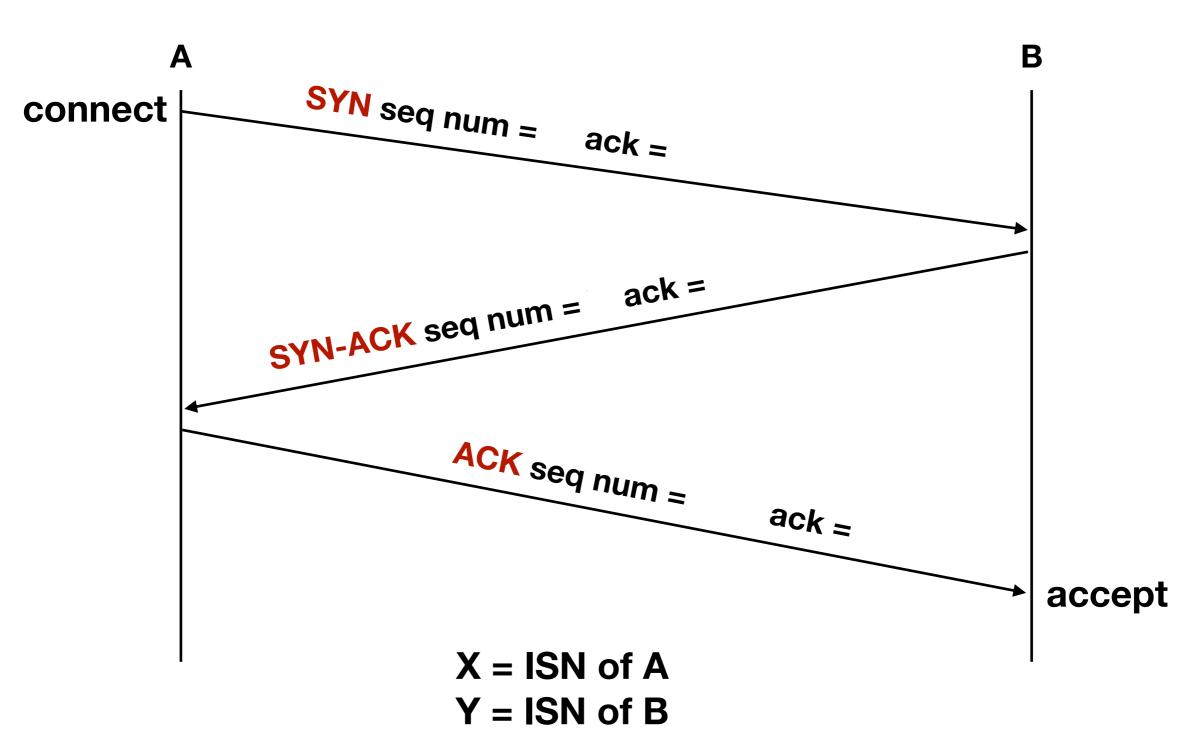
Practical Issue 2: TCP sequence prediction attack

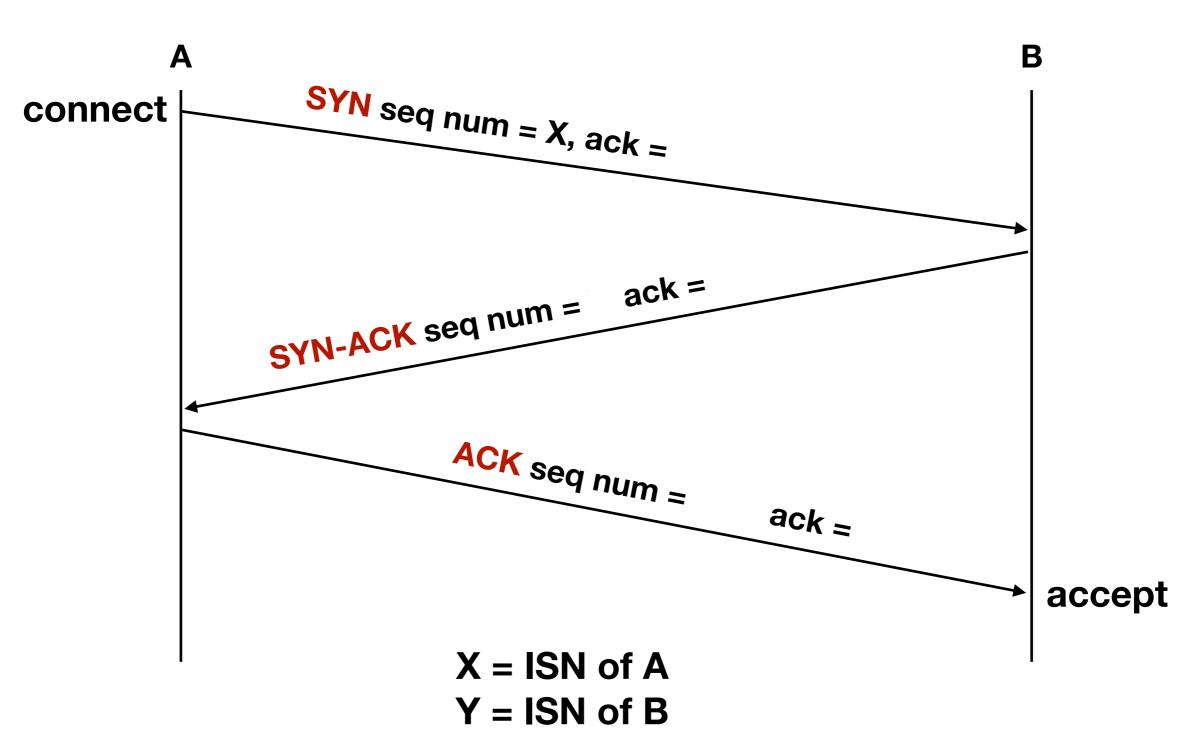
- Easy for attacker to guess a valid sequence number if ISN is always 0
- Can send counterfeit packets to receiver


TCP Connection Set up





- Host A sends a SYN ("Synchronize") packet to B
- Host B returns a SYN acknowledgement packet (SYN-ACK)
- Host A sends an ACK packet to acknowledge the SYN-ACK
- What should be the seq and ack numbers for these packets?

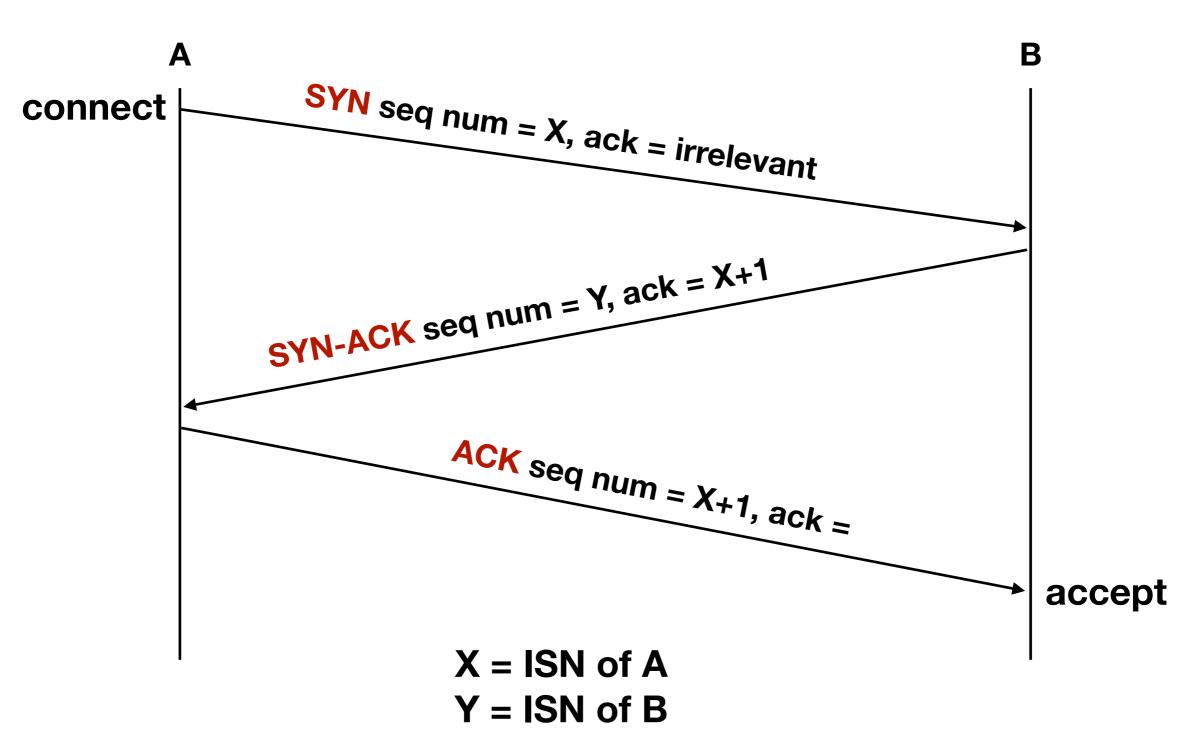


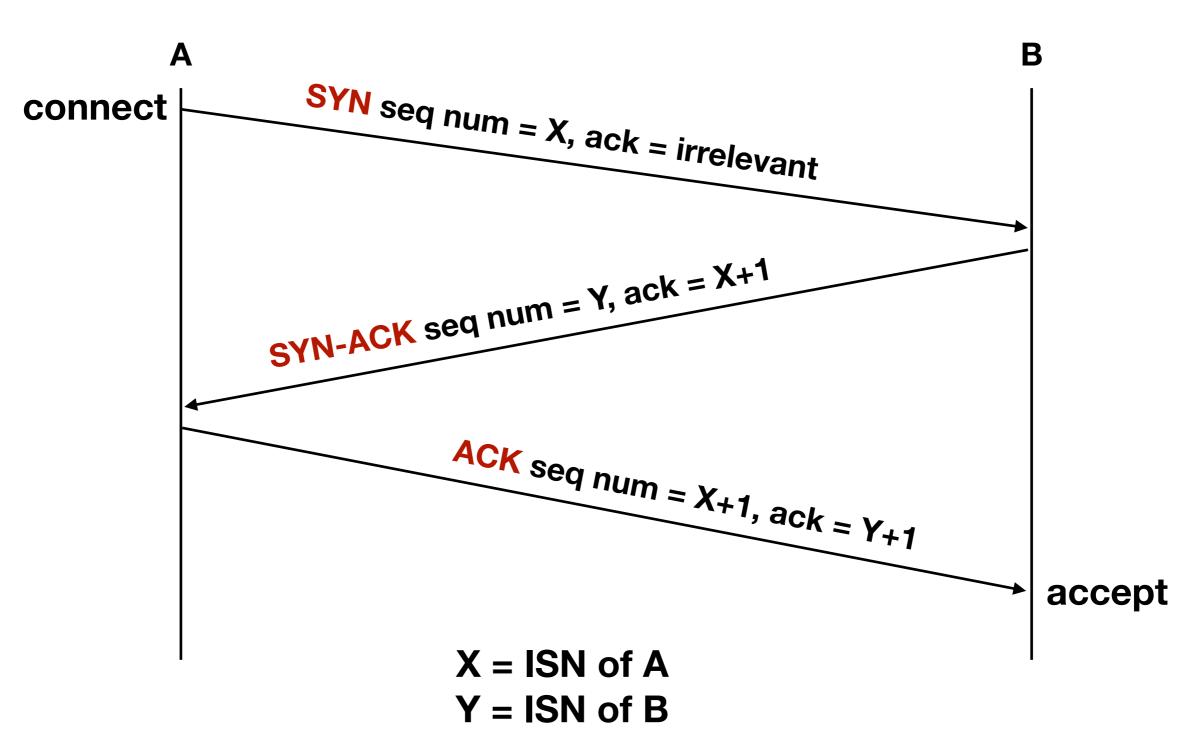








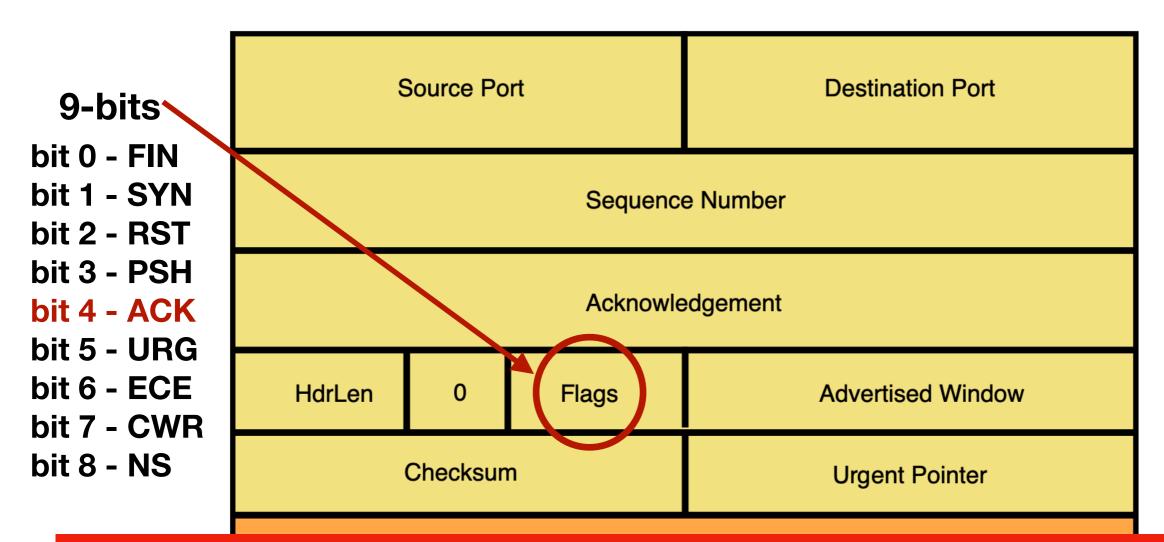




How to recognize Special Packets?

- How does TCP distinguish special packets (e.g., SYN, SYN-ACK) from normal data packets?
 - Answer: Using TCP Flags

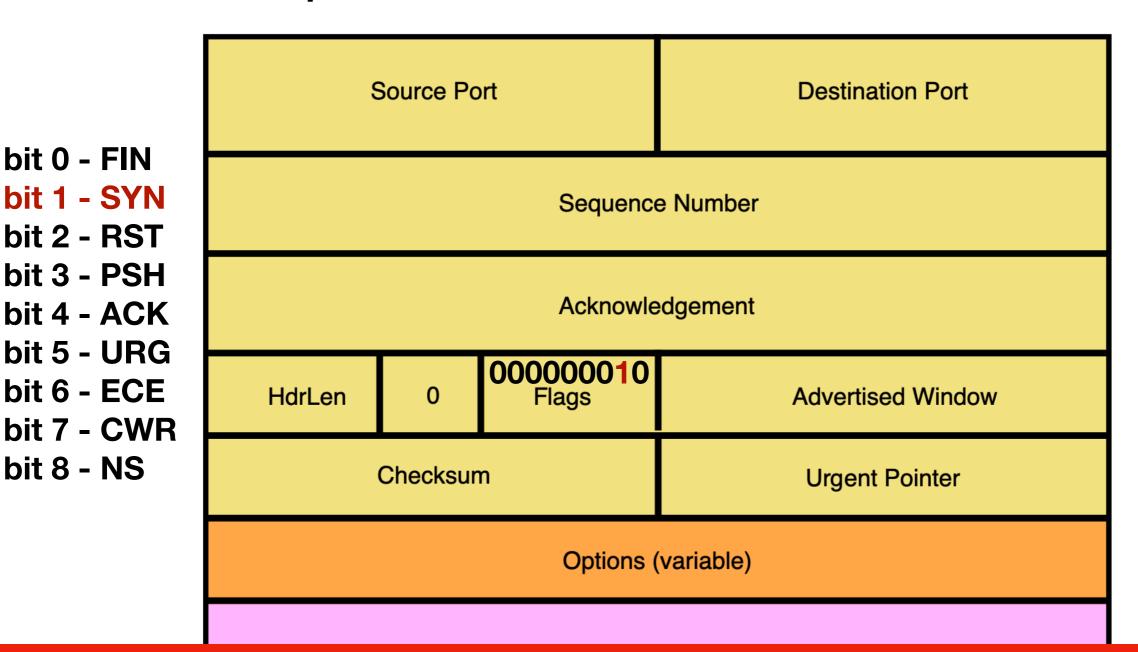
TCP Flags



IMPORTANT: The ACK flag bit is always set if the header contains a valid acknowledgement number


All TCP packets, except SYN packet and RST-REPLY packet, have ACK flag bit set!

Step 1: SYN Packet from A



A tells B it wants to open a connection NOTE: ACK flag bit is not set as no valid ack number


Step 2: SYN-ACK Packet from B



bit 8 - NS

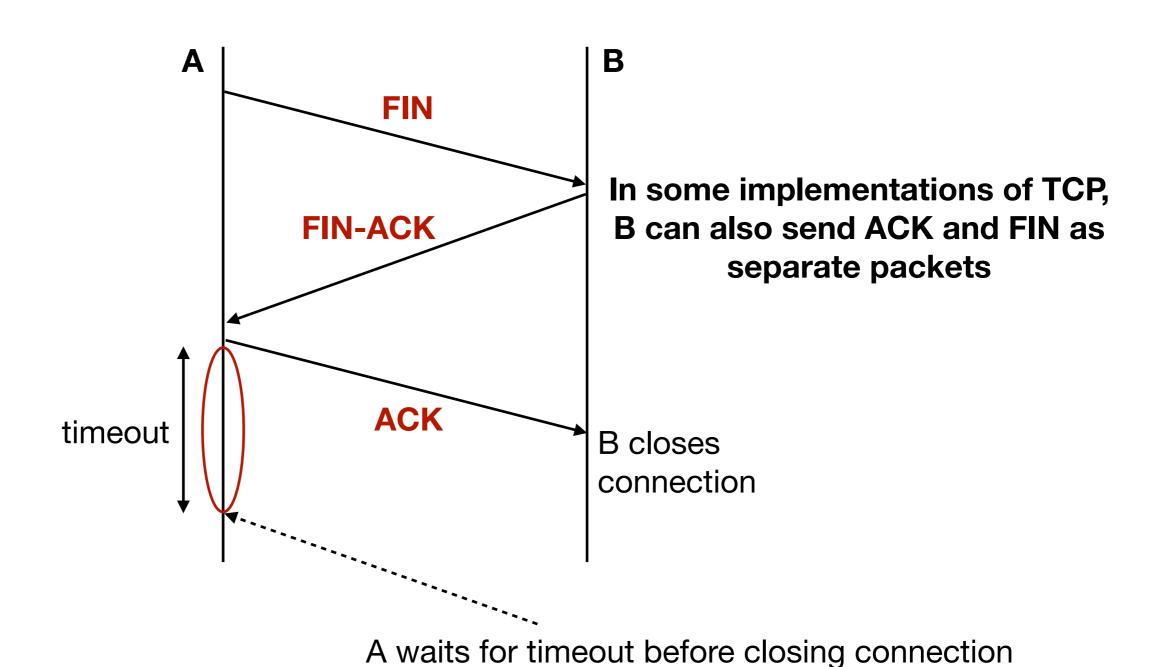
B tells A that it is ready to receive next byte on receiving this packet, A can start sending data

Step 3: ACK Packet from A



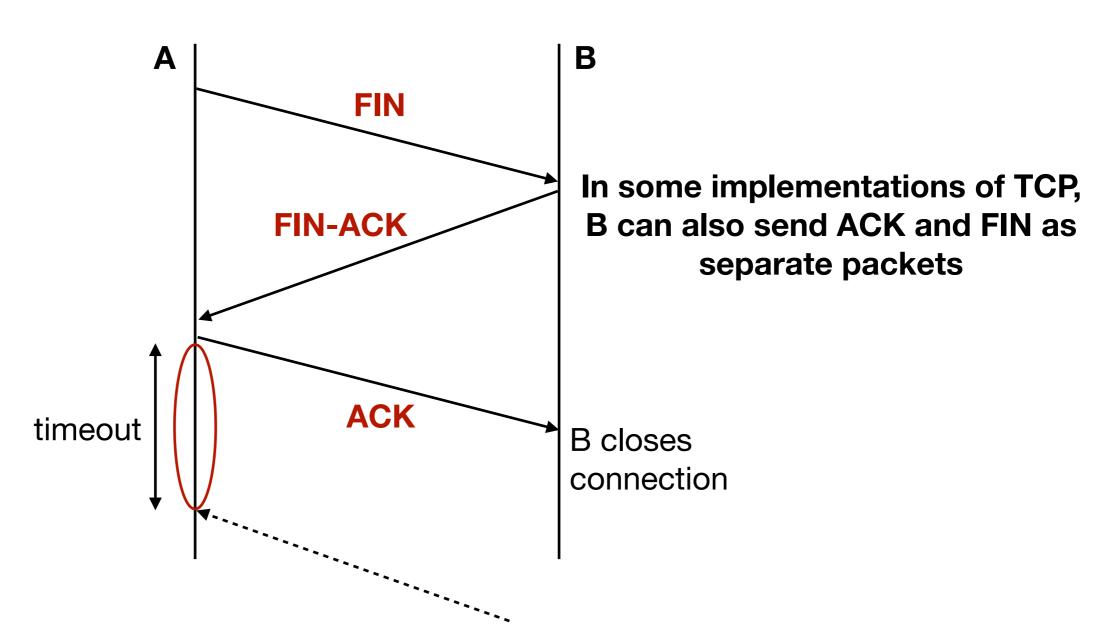
A tells B that it is also ready to receive next byte ... on receiving this packet, B can start sending data

TCP (Normal) Connection Termination



19

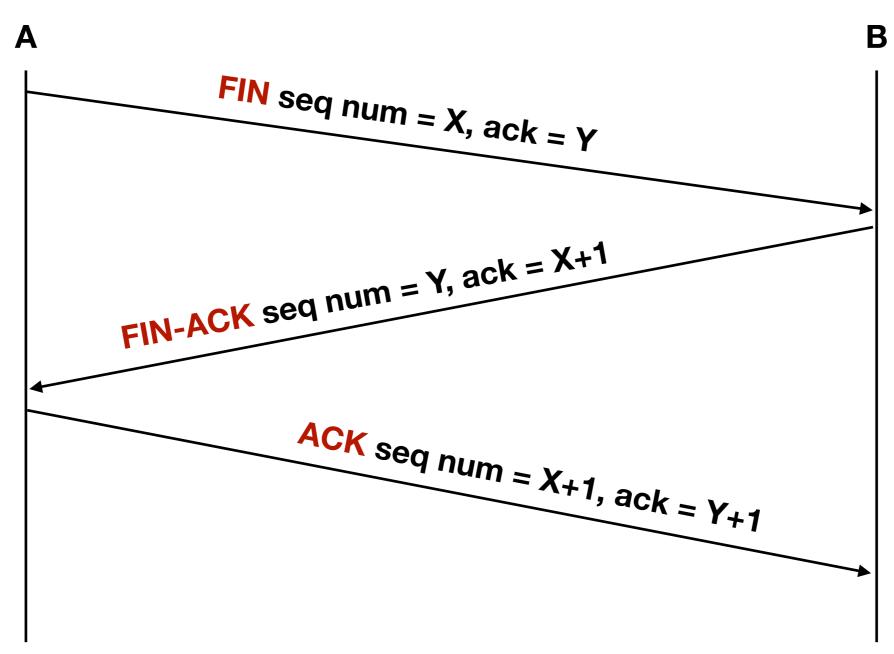
TCP (Normal) Connection Termination



A waits for timeout before closing connection

To make sure ACK is received by B, else A will receive a retransmitted FIN-ACK from B before timeout

3-Way Termination



X = A's current sequence number Y = sequence number of the next byte expected from B

Step 1: FIN Packet from A

bit 0 - FIN

bit 1 - SYN

bit 2 - RST

bit 3 - PSH

bit 4 - ACK

bit 5 - URG

bit 6 - ECE

bit 7 - CWR

bit 8 - NS

Source Port **Destination Port** Sequence Number Acknowledgement 000010001 HdrLen 0 Flags **Advertised Window** Checksum **Urgent Pointer** Options (variable) Data

Step 2: FIN-ACK Packet from B

bit 0 - FIN

bit 1 - SYN

bit 2 - RST

bit 3 - PSH

bit 4 - ACK

bit 5 - URG

bit 6 - ECE

bit 7 - CWR

bit 8 - NS

Source Port **Destination Port** Sequence Number Acknowledgement 000010001 HdrLen 0 Flags **Advertised Window** Checksum **Urgent Pointer** Options (variable) Data

Step 3: ACK Packet from A

bit 0 - FIN

bit 1 - SYN

bit 2 - RST

bit 3 - PSH

bit 4 - ACK

bit 5 - URG

bit 6 - ECE

bit 7 - CWR

bit 8 - NS

Source Port **Destination Port** Sequence Number Acknowledgement 000010000 HdrLen Flags 0 **Advertised Window** Checksum **Urgent Pointer** Options (variable) Data

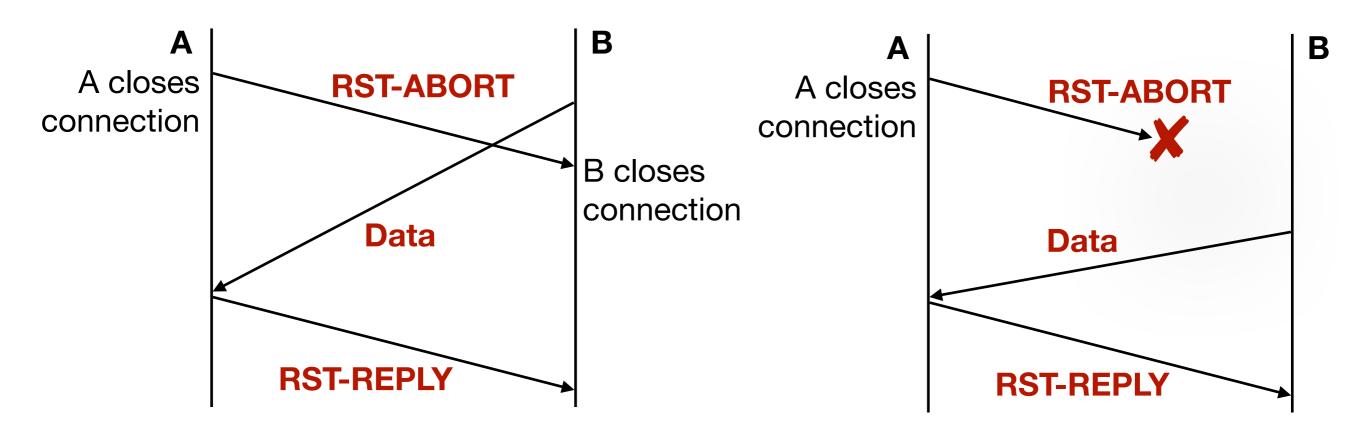
TCP (Abrupt) Connection Termination

A TCP endpoint can close a connection any time by sending a RESET (RST) pkt

\$ error: connection reset by peer

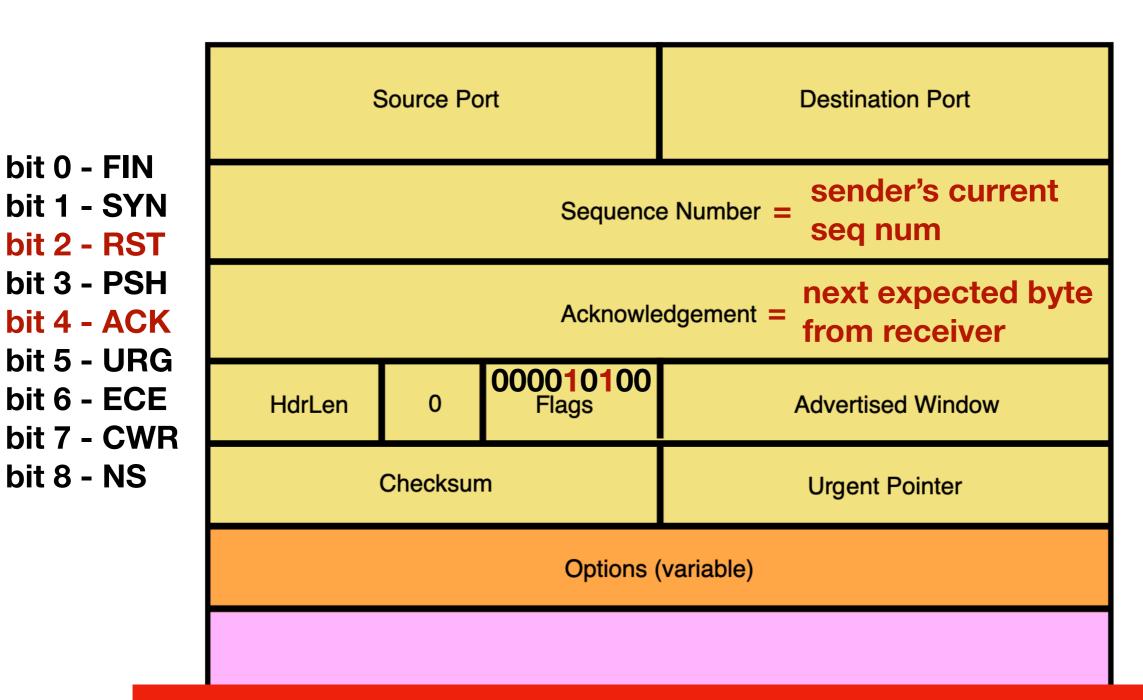
- Two kinds of RST packets:
 - 1. RST-ABORT: Generated when a connection is explicitly aborted by an endpoint (e.g., the socket at the endpoint is being killed)
 - 2. RST-REPLY: Generated on receipt of certain kinds of invalid packets (e.g., data packet for a connection that does not exist anymore)

TCP (Abrupt) Connection Termination



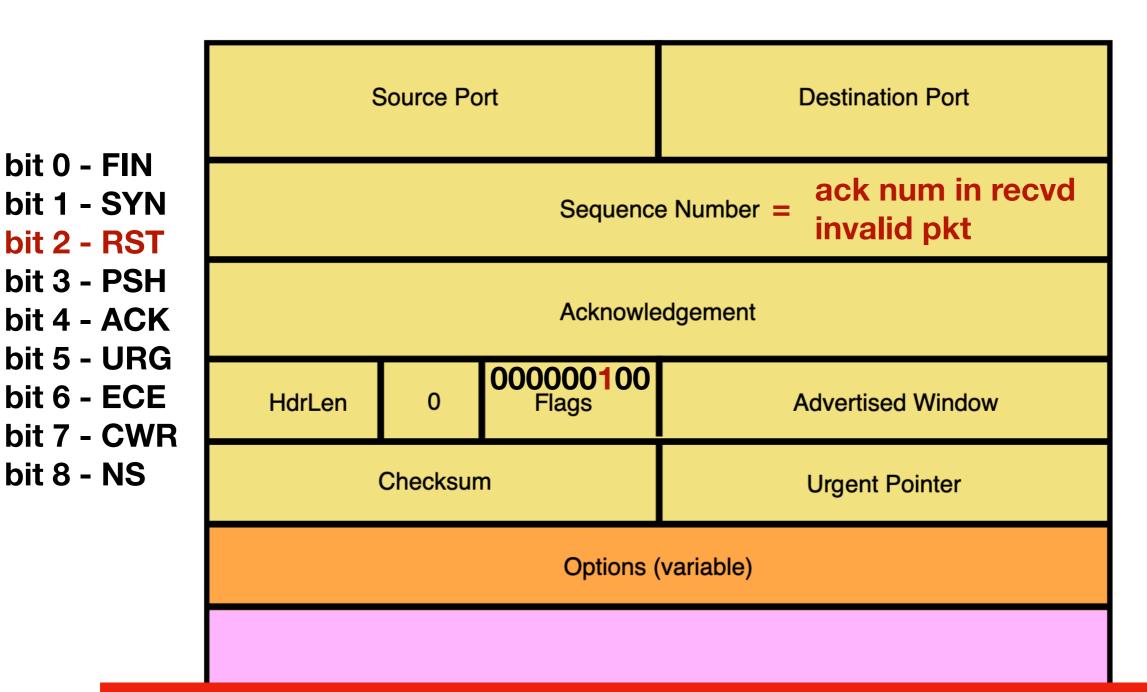
- A sends a RESET (RST-ABORT) packet to B to abruptly close connection
 - That's it... B closes its connection on receipt and does not ACK the RST packet
 - Thus, RST packet is not delivered reliably!
- If an endpoint receives data for a closed conn, it will trigger RST-REPLY from A

RST-ABORT Packet




ACK flag bit is set in RST-ABORT packets

RST-REPLY Packet



ACK flag bit is not set in RST-REPLY packets

TCP Full Timing Diagram



Connection set up

TCP is a "duplex" connection

- Data can be sent in both directions at the same time
- Data packets carry both data and ACK
 - Carry own data, ACK other endpoint's data
 - If ACK flag bit is set, ACK number in header is valid
 - All TCP data packets have ACK flag bit set

Connection termination

Data Packet

bit 0 - FIN

bit 1 - SYN

bit 2 - RST

bit 3 - PSH

bit 4 - ACK

bit 5 - URG

bit 6 - ECE

bit 7 - CWR

bit 8 - NS

Source Port **Destination Port** Sequence Number seq num of own data Acknowledgement ACK other endpoint's data 000010000 HdrLen 0 Flags **Advertised Window** Checksum **Urgent Pointer** Options (variable) Data carry own data

Remaining TCP Flags

- bit 0 FIN
- bit 1 SYN
- bit 2 RST
- bit 3 PSH
- bit 4 ACK
- bit 5 URG
- bit 6 ECE
- bit 7 CWR
- **bit 8 NS**

PSH Flag

bit 0 - FIN

bit 1 - SYN

bit 2 - RST

bit 3 - PSH


bit 4 - ACK

bit 5 - URG

bit 6 - ECE

bit 7 - CWR

bit 8 - NS



PSH Flag

- TCP by default can do data "batching" on both send and recv side
 - Send: TCP can batch data in send buffer into a single packet
 - Reduces header overhead and # of TCP packets sent to the network
 - Recv: TCP can wait to batch data in recv buffer before sending to app
 - Reduces number of receive calls from the application
- But, batching also adds delay ... PSH flag bit is used to disable batching
- Sender application can set PSH option in its socket interface
 - TCP sends out whatever data is present in the send buffer immediately
 - Also sets the PSH flag bit in the outgoing packet
- On the recv side, if TCP receives a packet with PSH flag bit set ...
 - ... does not wait to batch that data before sending to the application
- Desirable for interactive applications (e.g., chat applications)

URG Flag

bit 0 - FIN

bit 1 - SYN

bit 2 - RST

bit 3 - PSH

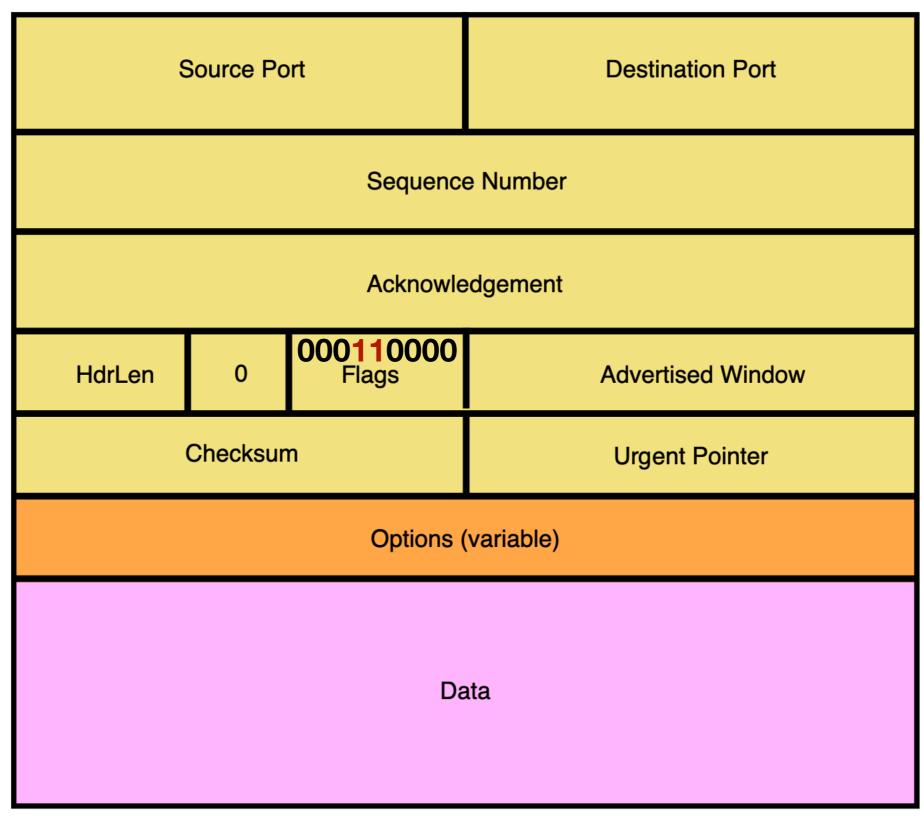
bit 4 - ACK

bit 5 - URG

bit 6 - ECE

bit 7 - CWR

bit 8 - NS



URG Flag

- TCP by default delivers data to application layer "in-order"
 - i.e., received data from network layer is added to the recv buffer ...
 - ... and dequeued and sent to application in a First-In-First-Out (FIFO) manner
- URGENT (URG) flag bit is used to deliver data "out-of-order"
 - If URG flag bit is set in a TCP packet,
 - Urgent data in TCP segment is delivered to application immediately
 - ... bypassing the recv buffer
 - Urgent Pointer in TCP header stores offset to the last urgent byte in TCP segment, starting from the first byte in the segment

Urgent Pointer

Source Port **Destination Port** bit 0 - FIN bit 1 - SYN Sequence Number bit 2 - RST bit 3 - PSH Acknowledgement bit 4 - ACK bit 5 - URG **000110000** Flags bit 6 - ECE HdrLen **Advertised Window** 0 bit 7 - CWR **bit 8 - NS** Checksum Urgent Pointer 16 bits Options (variable) **Urgent Data** Data

ECE, CWR, and NS Flags

```
bit 1 - SYN
bit 2 - RST
bit 3 - PSH
bit 4 - ACK
bit 5 - URG
bit 6 - ECE (ECN-Echo)
bit 7 - CWR (Congestion Window Reduced)
bit 8 - NS (Nonce Sum) Experimental
```

bit 0 - FIN

Transmission Control Protocol (TCP)

Reliable, In-order Delivery to Applications

- Application Multiplexing
- ACKs and Retransmissions
- Flow Control
- Congestion Control

Byte Stream Abstraction

Communicates with the applications using a byte stream and <u>not</u> packets

Connection-oriented

- Pairwise < sender, receiver > connection is established before sending data
- Used to maintain connection-specific state, e.g., initial seq num, window scaling factor, window size, etc.

Questions?